

# Air Quality (PM2.5) Forecasting

(3 horizons: +24h, +48h, +72h)

## 1. Objective

The goal of this project was to build an **end-to-end hourly pipeline** that:

- Pulls the latest air-quality and weather data for **Karachi**,
- Features (lags, rolling stats, calendar info),
- Trains and versions models for **24h, 48h, 72h ahead PM2.5**,
- Exposes the forecasts in a **Streamlit UI**
- Pushes the engineered features to **Hopsworks Feature Store** from GitHub Actions.

## 2. Data Ingestion (src/ingest.py)

**Data sources used:**

- **Open-Meteo Geocoding API**, to get latitude, longitude and timezone for the city (Karachi).
- **Open-Meteo Weather API**, to get hourly weather for the date range.
- **OpenAQ**, to get **hourly PM2.5** near Karachi. I searched for nearby locations/sensors within 25 km and aggregated hourly values.
- **Fallback**: if OpenAQ gives nothing, I call **Open-Meteo Air Quality (pm2\_5)** so the pipeline still works.

### **What the script does:**

- Pick a lookback window (in GitHub Action I used 90 days).
- Download weather hourly.
- Download PM2.5 hourly.
- Merge the two on timestamp and save to parquet files.

## 3. Feature Engineering (src/features.py)

### **Steps:**

- **Time features**
  - Hour, day-of-week, day-of-month, month
  - Reason: AQI often follows human/traffic and daily patterns.
- **Lag / rolling features on PM2.5**
  - pm25\_lag1 (1 hour ago)
  - pm25\_lag24 (same hour previous day)
  - pm25\_ma6 (6-hour moving average)
  - pm25\_ma24 (24-hour moving average)
  - pm25\_chg1 (current – previous hour)
  - Reason: AQI is very auto correlated; recent history is the strongest signal.
- **Future targets**
  - pm25\_tplus\_24
  - pm25\_tplus\_48
  - pm25\_tplus\_72
- **Save final training file to:**  
data/features/features.parquet

## 4. Model Training (src/train.py)

- Predict PM2.5 +24h
- Predict PM2.5 +48h
- Predict PM2.5 +72h

## Models compared per horizon:

- **Ridge Regression**, fast, good baseline
- **RandomForestRegressor**, non-linear, handles interactions
- **GradientBoostingRegressor**, added to improve expressiveness

For each horizon:

- Loaded data/features/features.parquet
- Did a train/test split
- Trained the 3 models
- Evaluated with **RMSE, MAE, R<sup>2</sup>**
- Picked the model with the **lowest RMSE**

## 5. Automation (GitHub Actions)

- **features-hourly.yml**
  - Every hour: checkout → install → run ingest.py → run features.py → commit updated features.
  - This keeps features.parquet fresh.
- **train-daily.yml**
  - Daily: checkout → install → run train.py
  - Copy the latest trained models into models/latest/
  - Hopsworks secrets exist → switch to Python 3.10 → install only HSFS → run src/push\_features\_hopsworks.py
  - Upload models as artifact

## 6. Streamlit UI (app/streamlit\_app.py)

- **Loads the latest trained models** from models/latest/
- **Gets the next hours' weather** from Open-Meteo (live)
- **Rebuilds the same feature row** as training (lags from your local pm25 parquet + weather) and runs prediction for 24/48/72h.

It displays:

- 3 big metrics: “PM2.5 forecast (+24h)” etc.
- color badge for AQI category (Good / Moderate / ...)
- a small SHAP-style explanation if available
- debug expanders for recent PM2.5 and training report

## 7. Hopsworks Integration

We **successfully pushed** the engineered features to Hopsworks Feature Store by installing only **hsfs[python]** on **Python 3.10 in CI**.

**UI:**

**AQI Forecast — 24h / 48h / 72h**

City: Karachi — times shown in Asia/Karachi (PKT)

Trained model winners (latest run):

```
["+24h": "ridge", "+48h": "rf", "+72h": "rf"]
```

Prediction reference time: 2025-11-07 05:00 PKT

PM2.5 forecast (+24h)	PM2.5 forecast (+48h)	PM2.5 forecast (+72h)
36.7 $\mu\text{g}/\text{m}^3$	36.6 $\mu\text{g}/\text{m}^3$	35.9 $\mu\text{g}/\text{m}^3$

Explain which forecast?

+24h

Why the +24h prediction?

feature	influence
dom	3.0297
pm25	2.0765
temperature_2m	0.6754
pm25_lag24	0.2584
pm25_lag1	0.2569
pm25_chg1	0.239
pm25_ma6	0.1913

```

▼ Training report (JSON)
{
  "version": "2025-11-07_1943",
  "feature_names": [
    0: "pm25",
    1: "pm25_lag1",
    2: "pm25_ma6",
    3: "pm25_ma24",
    4: "pm25_chg1",
    5: "temperature_2m",
    6: "relative_humidity_2m",
    7: "wind_speed_10m",
    8: "surface_pressure",
    9: "hour",
    10: "dow",
    11: "dom",
    12: "month"
  ],
  "horizons": {
    "h24": {
      "best_model": "ridge",
      "metrics": {
        "ridge": {
          "rmse": 14.485624125337,
          "mae": 10.187771451389758,
          "r2": 0.31532355556542713
        }
      },
      "rf": {
        "rmse": 17.717570918166017,
        "mae": 11.785370577281185,
        "r2": -0.022785476110592784
      }
    }
  }
}

```

## FEATURES:

```

def add_time_features(df: pd.DataFrame):
    df = df.copy()
    df[TIME_COL] = pd.to_datetime(df[TIME_COL], utc=True)

    df["hour"] = df[TIME_COL].dt.hour
    df["dow"] = df[TIME_COL].dt.dayofweek
    df["dom"] = df[TIME_COL].dt.day
    df["month"] = df[TIME_COL].dt.month

    print("time featuring\n", df.head())
    return df

def add_lag_rolling_change(df: pd.DataFrame):
    """
    Add PM2.5 history columns:
    - pm25_lag1 (1 hour ago)
    - pm25_lag24 (24 hours ago)
    - pm25_ma6 (6-hour moving avg)
    - pm25_ma24 (24-hour moving avg)
    - pm25_chg1 (current - 1 hour ago)
    """
    df = df.copy()

    df[f"{TARGET_COL}_lag1"] = df[TARGET_COL].shift(1)
    df[f"{TARGET_COL}_lag24"] = df[TARGET_COL].shift(24)
    df[f"{TARGET_COL}_ma6"] = df[TARGET_COL].rolling(6).mean()
    df[f"{TARGET_COL}_ma24"] = df[TARGET_COL].rolling(24).mean()
    df[f"{TARGET_COL}_chg1"] = df[TARGET_COL] - df[f"{TARGET_COL}_lag1"]

    print("Lag , Rolling\n", df.head())
    return df

def add_future_targets(df: pd.DataFrame):
    df = df.copy()

    for h in [24, 48, 72]:
        col_name = f"{TARGET_COL}_tplus_{h}"
        df[col_name] = df[TARGET_COL].shift(-h)

    print("Future Targets\n", df.head())
    return df

```